

```
In [9]: %load_ext autoreload
%autoreload 2

# Control figure size
figsize=(15, 4)

from sklearn.neighbors import KernelDensity
from sklearn.model_selection import GridSearchCV
from util import util
import numpy as np
from matplotlib import pyplot as plt
import pandas as pd
from sklearn.gaussian_process.kernels import RBF, ConstantKernel, WhiteKernel
from sklearn.gaussian_process.kernels import ExpSineSquared, DotProduct, Mat
from sklearn.gaussian_process import GaussianProcessRegressor
from sklearn.metrics import mean_squared_error
import os
import warnings
warnings.filterwarnings(action='once')
warnings.filterwarnings('ignore')

# Load data
data_folder = os.path.join('.', 'data', 'nab')
file_name = os.path.join('realTraffic', 'occupancy_6005.csv')
data, labels, windows = util.load_series(file_name, data_folder)
data_d = ddata = data.resample('5min').mean()
```

The autoreload extension is already loaded. To reload it, use:
 %reload_ext autoreload


Filling Missing Values in Traffic Data

Back to the Traffic Data

We will try to use a Gaussian Process on our new traffic data

We will start directly from the dense series:

```
In [10]: util.plot_series(data_d, show_markers=True, figsize=figsize)
```



- There is a *period*, *non-zero mean*, *local oscillations*. And of course a *huge gap*

Time and Period

The input of the GP is going to be the *time* of the observation

Unfortunately, the GP in scikit-learn cannot handle `DateTime` objects

- Therefore, we will convert all time steps in numeric format
- We will use a simple time equivalent, namely 1 step = 1 time unit

```
In [11]: data_dt = data_d.copy()
data_dt['time'] = np.arange(len(data_dt))
```

Before we start, it would be very useful to estimate the period

Due to the missing values, the series has a *non-uniform sampling frequency*

- Autocorrelation plots cannot be used in this case
- The standard FFT is also not applicable (we could use a [non-uniform FFT](#))

Thankfully, this is traffic data! So we can bet there is a *weekly period*

Process Outline

We have no ground truth: how are we going to evaluate the kernels?

We will use the same trick we used before:

- We will *focus on a portion* of our sequence
 - ...One with relatively few missing values
- Then we will *artificially remove* part of the data points
 - This will form the ground truth for our evaluation

Main idea: use *part of our data as a validation set*


Which quality metric?

- Thanks to the availability of confidence intervals...
- ...We can compute the *likelihood* of our validation set!
- Using the MSE would do the same, only with more assumptions

Training and Validation Data

We will use for training (and validation) this stretch of the series

```
In [12]: segment = data_dt[(data_dt.index >= '2015-09-09') & (data_dt.index < '2015-09-17')]
util.plot_series(segment['value'], show_markers=True, figsize=figsize)
```



- We made sure to include *at least one full week*

Training and Validation Data

We can now separate training and validation data:

```
In [14]: tmp = segment.dropna()

np.random.seed(42)
idx = np.arange(len(tmp))
np.random.shuffle(idx[1:-1]) # do not shuffle the first/last point
t = idx[1]; idx[1] = idx[-1]; idx[-1] = t # keep first/last points in the le

sep = 2*len(idx) // 3
trdata = tmp.iloc[idx[:sep]]
tsdata = tmp.iloc[idx[sep:]]
```

- We are keeping 2/3 of the data for training
- Since we are using the dense series, we need to discard NaNs (`dropna`)
- Since we are doing interpolation...
- ...It's a good idea to keep the first and last point in the training set

A Starting Kernel

Let's try with a relatively simple kernel

```
In [33]: kernel = WhiteKernel(1e-3, (1e-4, 1e1))
#kernel = WhiteKernel(1e-3, (1e-4, 1e-1))
kernel += ConstantKernel(1, (1e-2, 1e2)) * RBF(1, (1e-1, 1e1))
```

```

np.random.seed(42)
gp = GaussianProcessRegressor(kernel=kernel, n_restarts_optimizer=3)
gp.fit(trdata[['time']], trdata['value'])
print(gp.kernel_)

```

WhiteKernel(noise_level=3.76) + 3.61**2 * RBF(length_scale=10)

Then we obtain the predictions:

```

In [34]: pred, std = gp.predict(segment[['time']], return_std=True)
pred = pd.Series(index=segment.index, data=pred)
std = pd.Series(index=segment.index, data=std)

```


A Starting Kernel

Let's have a look at the predictions

```

In [35]: util.plot_gp(target_samples=tsdata['value'], samples=trdata['value'], pred=p

```



- Not so bad, but not so good either
- We have *very wide confidence intervals!*

A Starting Kernel

Let's compute the (log) likelihood of the validation data

```

In [36]: from scipy.stats import norm

# Obtain predictions for the validation data
pred_ts = pred[tsdata.index]
std_ts = std[tsdata.index]

ldens = norm.logpdf(tsdata['value'], pred_ts, std_ts)
ll = np.sum(ldens)
print(f'Log likelihood of the validation set: {ll:.2f}')

```

Log likelihood of the validation set: -1123.17

- This is our reference value

- We will try to beat it by *improving the kernel*

A Second Kernel

Let's add *the period*

We can choose the bounds so as to focus on a weekly period

```
In [37]: kernel = WhiteKernel(1e-3, (1e-4, 1e1))
#kernel = WhiteKernel(1e-3, (1e-4, 1e-1))
kernel += ConstantKernel(1, (1e-2, 1e2)) * RBF(1, (1e-1, 1e1))
kernel += ConstantKernel(1, (1e-2, 1e2)) * ExpSineSquared(1, 2000, (1e-1, 1e1))

np.random.seed(42)
gp = GaussianProcessRegressor(kernel=kernel, n_restarts_optimizer=3)
gp.fit(trdata[['time']], trdata['value'])
print(gp.kernel_)
```

```
WhiteKernel(noise_level=0.56) + 1.88**2 * RBF(length_scale=0.1) + 5.37**2 *
ExpSineSquared(length_scale=0.1, periodicity=2.01e+03)
```

Then we obtain the new predictions:


```
In [38]: pred, std = gp.predict(segment[['time']], return_std=True)
pred = pd.Series(index=segment.index, data=pred)
std = pd.Series(index=segment.index, data=std)
```

A Second Kernel

Both predictions and likelihood are now better

```
In [39]: util.plot_gp(target_samples=tsdata['value'], samples=trdata['value'], pred=pred,
ldens = norm.logpdf(tsdata['value'], pred[tsdata.index], std[tsdata.index])
print(f'Log likelihood of the validation set: {np.sum(ldens):.2f}')
```

Log likelihood of the validation set: -1110.69



- ...But the confidence intervals are still large

Transferring the Results

We now need to obtain predictions *for the whole series*

We would prefer to *avoid training again the kernel* parameters

- The large number of missing value may be problematic
- ...And the training time would be very large

...But we *also* really wish to use *all available observations*

- ...Not just those considered when training the kernel

With Gaussian Processes, we can do both

There is *no need to train again the kernel* every time new observations arrive

- We can build a new Σ matrix using *the new observations* and *the old kernel*

Transferring the Results

We reuse the kernel by passing it as argument when building a new GP:

```
In [40]: gp2 = GaussianProcessRegressor(kernel=gp.kernel_, optimizer=None)
```

Passing `optimizer=None` will *disable optimization* at training time

So that calling `fit` will just take into account a new set of observations

```
In [41]: tmp = data_dt.dropna() # The whole series (NaNs excluded)
gp2.fit(tmp[['time']], tmp['value']);
```


Now we can obtain predictions for the whole series:

```
In [42]: pred_all, std_all = gp2.predict(data_dt[['time']], return_std=True)
pred_all = pd.Series(index=data_dt.index, data=pred_all)
std_all = pd.Series(index=data_dt.index, data=std_all)
```

Transferring the Results

Let's have a look at all the predictions

```
In [43]: util.plot_gp(samples=data_dt['value'], pred=pred_all, std=std_all, figsize=f
```



- We actually managed to (partially) reconstruct *even the large gap!*
-But we *still* have those large confidence intervals

Transferring the Results

The confidence intervals are still very large

- This is in part understandable, give the presence of wide variations
- ...But at least one point is a bit strange

The confidence intervals are large *even for the night hours!*

There are two reasons:

- There are *fewer samples* at nighttime
 - As we get far from the samples the confidence drops (quickly, in our case)
- No traditional GP kernel can represent *input-dependent variance*
 - All kernels are about covariance, not variance
 - The lone exception is the `WhiteKernel`, which is not input dependent

Can we deal with this issues?

Standard Deviation Model

Multiplicative Ensemble

We can deal with the input-dependent variance *in a separate model*

We are going to build an *ensemble* of predictive models

- Classical ensembles (e.g. Random Forests, Gradient Boosting) are based on *sums*
- But that's not going to work with variance, since:

$$\text{Var}(x + \alpha) = \text{Var}(x)$$

However, variance can be *scaled via multiplication*:

In particular:

$$\text{Var}(\alpha x) = \alpha^2 \text{Var}(x)$$

- So we can use a "multiplicative" ensemble

Multiplicative Ensemble

Our model will become *the product of two models*

Formally, we will have:

$$g(x, \lambda) f(x, \theta)$$

- f , with parameters θ will be a Gaussian Process
- g , with parameters λ will be our *variance model* (or *standard deviation model*)

On the training set, we wish to have:

$$g(x_i, \lambda) f(x_i, \theta) \simeq y_i \quad \Rightarrow \quad f(x_i, \theta) \simeq \frac{y_i}{g(x_i, \lambda)}$$

- The Gaussian Process will need to learn a series with a variance *altered by g*
- The variance of each point y_i will be divided by $g(x_i, \lambda)^2$

Standard Deviation Model

We now need to choose our variance model g

- Since we have discrete time and a natural period (a week)
- ...We could use a simple map (time of the week $\rightarrow$ standard deviation)

Let's add a "hour of the week" information to our data:

The chosen time unit is actually irrelevant

```
In [44]: data_dtw = data_dt.copy()
data_dtw['how'] = 24 * data_dt.index.weekday + data_dt.index.hour + data_dt.index.minute
data_dtw.head()
```


Out [44]:

	value	time	how
--	-------	------	-----

timestamp			
2015-09-01 13:45:00	3.06	0	37.750000
2015-09-01 13:50:00	6.44	1	37.833333
2015-09-01 13:55:00	5.17	2	37.916667
2015-09-01 14:00:00	3.83	3	38.000000
2015-09-01 14:05:00	4.50	4	38.083333

Standard Deviation Model

Then we can compute the standard deviation via a pandas `groupby` operation:

```
In [45]: how_std = data_dtw.groupby('how').agg({'value': ['std', 'count']})
```

- `agg` allows to apply multiple aggregation functions to multiple columns

The resulting table has a *hierarchical column index*

Let's see some statistics about the value counts:


```
In [46]: how_std['value', 'count'].describe() # Notice the use of two names
```

```
Out [46]: count      2016.000000
mean          1.177083
std           0.850567
min           0.000000
25%           1.000000
50%           1.000000
75%           2.000000
max           3.000000
Name: (value, count), dtype: float64
```

Standard Deviation Model

Let's have a look at the *standard deviation* values

```
In [47]: util.plot_series(how_std['value', 'std'], show_markers=True, figsize=figsize)
```



- There are *many missing values* (as expected from the counts)!
- Note: this is a time-to-stdev map, not our original series!

Choosing the Correct Interval

There are too many gaps in our data to compute σ with this granularity

There are a few possible solutions:

- *Filling the gaps* also in the standard deviation data
- ...Or using a *coarser time unit*
- ...Or choosing a *shorter period* (e.g. one day)

Let's try using **hour-long intervals (rather than 5 minutes)**


```
In [48]: data_dtw = data_dt.copy()
data_dtw['how'] = np.round(24 * data_dt.index.weekday + data_dt.index.hour + 1)
how_std = data_dtw.groupby('how').agg({'value': ['std', 'count']})
```

- We need to check that *no std value is missing*
- ...But also that *the counts are large enough* for a reliable computation

Choosing the Correct Interval

Let's look again at the standard deviation values

```
In [49]: util.plot_series(how_std['value', 'std'], show_markers=True, figsize=figsize)
```




- There are *still a few gaps*

Choosing the Correct Interval

...And let's see the value counts

```
In [50]: util.plotBars(how_std['value', 'count'], figsize=figsize)
```



- Several value counts are too low for σ to be reliable

Choosing the Correct Interval

We can try again with two-hour intervals

```
In [51]: data_dtw = data_dt.copy()
data_dtw['how'] = 2*np.round(0.5*(24 * data_dt.index.weekday + data_dt.index
how_std = data_dtw.groupby('how').agg({'value': ['std', 'count']})
```

Let's check some information about the value counts:

- Ideally, we wish them at ~ 30


```
In [52]: how_std['value', 'count'].describe()
```

```
Out[52]: count      85.000000
mean       27.917647
std        16.132791
min         3.000000
25%        16.000000
50%        22.000000
75%        42.000000
max        63.000000
Name: (value, count), dtype: float64
```

Choosing the Correct Interval

Let's look the standard deviation with two-hour intervals

```
In [53]: util.plot_series(how_std['value', 'std'], show_markers=True, figsize=figsize)
```



- Finally, *no more missing values* and *decently large counts*

Upsampling and Smoothing

We managed to have reasonable standard deviation values..

...But our map/table has a *very* coarse time unit!

- Using it would lead to *sharp variations* in our predicted standard deviation

We will now proceed to mitigate the problem

We will start by *upsampling*, i.e. switching to a finer grain time unit:

```
In [54]: how_values = np.unique(24 * data_dt.index.weekday + data_dt.index.hour + dat
how_std_d = pd.DataFrame(index=sorted(how_values), columns=['std'], data=np.
how_std_d['std'] = how_std['value', 'std'])
```


This process leads to many missing values, that we fill via *linear interpolation*:

```
In [55]: how_std_di = how_std_d.interpolate(method='linear')
```

Upsampling and Smoothing

Let's see the oversampled series

```
In [56]: util.plot_series(how_std_di['std'], show_markers=True, figsize=figsize, use_
```



- The plot is the same as before, but there are *more values on the x-axis*

Upsampling and Smoothing

We will smooth the cuve via a simple low-pass filter

I.e. the *Exponentially Weighted Moving Average*

- This is a form of discrete filter, given by the recursion:

$$ss_i = \begin{cases} x_i & \text{if } i = 1 \\ \alpha x_i + (1 - \alpha)s_{i-1} & \text{otherwise} \end{cases}$$

- s_i is the i -th element of the output (smoothed) series
- α is called the *smoothing factor* and is equal to $1/(1 + \tau)$

Why not using a simple moving average?

- A moving average gives the same weight to all observations...
- ...whereas in our case "*recent*" observations are *more important*
- I.e. the stdev from the original table should still be the dominant value


Upsampling and Smoothing

In pandas, we can use the `ewm`) iterator, plus the `mean` aggregation function

`DataFrame.ewm(com=None, ...).mean()`

- The `com` parameter corresponds to τ

```
In [57]: how_std_ds = how_std_di[['std']].ewm(com=12).mean()
util.plot_series(how_std_ds['std'], show_markers=True, figsize=figsize, use_
```



- We chose τ = the number of time steps in one hour

Gaussian Process for the Ensemble

Transforming the Dataset

We can now learn the Gaussian Process for our Ensemble

For this, we need to *transform the original series* using the stdev model

- We start by augmenting our dataset with the "hour of the week information"
- ...Then, we associate each data point to the predicted standard deviation

```
In [58]: data_dt['how'] = 24 * data_dt.index.weekday + data_dt.index.hour + data_dt.index.minute
data_dts = data_dt.join(how_std_ds, on='how')
data_dts.head()
```

```
Out[58]:
```


	value	time	how	std
timestamp				
2015-09-01 13:45:00	3.06	0	37.750000	2.822412
2015-09-01 13:50:00	6.44	1	37.833333	2.862934
2015-09-01 13:55:00	5.17	2	37.916667	2.904604
2015-09-01 14:00:00	3.83	3	38.000000	2.947336
2015-09-01 14:05:00	4.50	4	38.083333	2.981964

- We relied on the `join` method from `pandas`

Transforming the Dataset

Now we can actually transform the series values

```
In [59]: data_dts['value_tr'] = data_dts['value'] / data_dts['std']
util.plot_series(data_dts['value_tr'], show_markers=True, figsize=figsize)
```



- The series *has changed considerably*: this is not a simple standardization

Training Data

We can now select a sub-sequence of the data for learning the kernel

...Which is necessary, since the data has changed!

```
In [60]: segment_s = data_dts[(data_dts.index >= '2015-09-09') & (data_dts.index < '2015-09-17')]
```

We separate training and validation data as we did before:

```
In [61]: tmp = segment_s.dropna()

np.random.seed(42)
idx = np.arange(len(tmp))
np.random.shuffle(idx[1:-1]) # no not shuffle the first/last point
t = idx[1]; idx[1] = idx[-1]; idx[-1] = t # keep first/last points in the le

sep = 2*len(idx) // 3
trdata_s = tmp.iloc[idx[:sep]]
tsdata_s = tmp.iloc[idx[sep:]]
```

Learning the Kernel Parameters

We can now learn the kernel parameters

We can use the same starting parameters (*priors*) as before:

```
In [63]: kernel_s = WhiteKernel(1e-3, (1e-4, 1e1))
#kernel_s = WhiteKernel(1e-3, (1e-4, 1e-1))
kernel_s += ConstantKernel(1, (1e-2, 1e2)) * RBF(1, (1e-1, 1e1))
kernel_s += ConstantKernel(1, (1e-2, 1e2)) * ExpSineSquared(1, 2000, (1e-1,

np.random.seed(42)
gp_s = GaussianProcessRegressor(kernel=kernel_s, n_restarts_optimizer=3)
```

```
gp_s.fit(trdata_s[['time']], trdata_s['value_tr'])
print(gp_s.kernel_)
```

WhiteKernel(noise_level=0.81) + 0.1**2 * RBF(length_scale=0.319) + 1.67**2 * ExpSineSquared(length_scale=0.1, periodicity=2.01e+03)


Then we obtain the predictions:

```
In [64]: pred_s, std_s = gp_s.predict(segment_s[['time']], return_std=True)
pred_s = pd.Series(index=segment_s.index, data=pred_s)
std_s = pd.Series(index=segment_s.index, data=std_s)
```

Learning the Kernel Parameters

Let's look at the predictions on the training data

```
In [65]: util.plot_gp(target_samples=tsdata_s['value_tr'], samples=trdata_s['value_tr'])
```



- For sake of simplicity, we will not try to improve the kernel

Transferring the Results

Now we obtain predictions for the missing values in the *transformed* series

Again, we reuse the kernel and add the observations:

```
In [66]: gp2_s = GaussianProcessRegressor(kernel=gp_s.kernel_, optimizer=None)
tmp_s = data_dts.dropna() # The whole series (NaNs excluded)
gp2_s.fit(tmp_s[['time']], tmp_s['value_tr']);
```

Then we can obtain predictions (and confidence intervals) for the whole series


```
In [67]: pred_all_s, std_all_s = gp2_s.predict(data_dts[['time']], return_std=True)
pred_all_s = pd.Series(index=data_dts.index, data=pred_all_s)
std_all_s = pd.Series(index=data_dts.index, data=std_all_s)
```

- Of course are still referring to the transformed series

Transferring the Results

Let's have a look at all the predictions

```
In [68]: util.plot_gp(samples=data_dts['value_tr'], pred=pred_all_s, std=std_all_s, f
```




- They are not so easy to interpret, since they refer to the transformed series

Predictions for the Original Series

We obtain predictions *for the original series* by injecting back the variance

```
In [70]: pred_all2_s = pd.Series(index=data_dts.index, data=pred_all_s * data_dts['std']  
std_all2_s = pd.Series(index=data_dts.index, data=std_all_s * data_dts['std']  
util.plot_gp(samples=data_dts['value'], pred=pred_all2_s, std=std_all2_s, fig
```




- Due to the properties of variance
- ...We can *just multiply also the standard deviation*

Predictions for the Original Series

For comparison, here are the results for the previous Gaussian Process

```
In [71]: util.plot_gp(samples=data_dt['value'], pred=pred_all, std=std_all, figsize=f
```



- The new confidence intervals are *much tighter*

Fill with Predictions and Samples

We can fill the missing values using the *predictions*

This will fill each missing value using the Maximum A Posteriori

```
In [72]: mask = data_d['value'].isnull() # we need to fill only the NaNs
data_filled_pred = data_d.copy()
data_filled_pred.loc[mask, 'value'] = np.maximum(0, pred_all2_s[mask])
```

But with GPs, we can also *sample from the distribution*


```
In [73]: tmp = data_dts[mask]
sample_ms = gp2_s.sample_y(tmp[['time']], random_state=42).ravel()
data_filled_samples = data_d.copy()
data_filled_samples.loc[mask, 'value'] = np.maximum(0, sample_ms * tmp['std']
```

- `sample_y` returns a matrix: we used `ravel` to have a single dimension
- In both cases, we *clip values at zero* (no less than 0 occupancy)

Filling with Predictions and Samples

Here's the series filled *using predictions*:

```
In [74]: util.plot_series(data_filled_pred, figsize=figsize)
```




- The use of MAPs is evident in the large gap

Filling with Predictions and Samples

Here's the series filled *using samples*:

```
In [75]: util.plot_series(data_filled_samples, figsize=figsize)
```



- There no evidently "fake" sections, now! ...Except for the effect of clipping at 0