

```
In [1]: # =====
# Notebook setup
# =====

%load_ext autoreload
%autoreload 2

# Control figure size
figsize=(14, 4)

from util import util
import os
import numpy as np
from tensorflow.keras import callbacks
import numpy as np
import pandas as pd
from matplotlib import pyplot as plt
from sklearn.gaussian_process.kernels import RBF, WhiteKernel, ConstantKernel
from sklearn.gaussian_process import GaussianProcessRegressor
from tensorflow import keras
from bayes_opt import BayesianOptimization

import warnings
warnings.simplefilter("ignore")

# Load data
data_folder = os.path.join '..', 'data'
data = util.load_data(data_folder)

# Identify input columns
dt_in = list(data.columns[3:-1]) # Exclude metadata

# Focus on a subset of the data
data_by_src = util.split_by_field(data, field='src')
dt = data_by_src['train_FD004']

# Split training and test machines
tr_ratio = 0.75
np.random.seed(42)
machines = dt.machine.unique()
np.random.shuffle(machines)

sep = int(tr_ratio * len(machines))
tr_mcn = machines[:sep]
ts_mcn = machines[sep:]

tr, ts = util.partition_by_machine(dt, tr_mcn)

# Standardization and normalization
trmean = tr[dt_in].mean()
trstd = tr[dt_in].std().replace(to_replace=0, value=1) # handle static fields
ts_s = ts.copy()
ts_s[dt_in] = (ts_s[dt_in] - trmean) / trstd
tr_s = tr.copy()
```

```

tr_s[dt_in] = (tr_s[dt_in] - trmean) / trstd

trmaxrul = tr['rul'].max()
ts_s['rul'] = ts['rul'] / trmaxrul
tr_s['rul'] = tr['rul'] / trmaxrul

# Define a cost model
failtimes = dt.groupby('machine')['cycle'].max()
safe_interval = failtimes.min()
maintenance_cost = failtimes.max()
cmodel = util.RULCostModel(maintenance_cost=maintenance_cost, safe_interval=

```

Bayesian (Surrogate-Based) Optimization

Bayesian Optimization

We will use an approach known as *Surrogate-Based Bayesian Optimization*

- It is designed to optimize *blackbox functions*
- I.e. functions with an unknown structure, that can only be evaluated

Formally, they address problems in the form:

$$\min_{x \in B} f(x)$$

- Where B is a box, i.e. a specification of bounds for each component of x
- In our case, the decision variable x would be θ
- ...And the function to be optimized would be the cost

The functions are typically assumed to be *expensive to evaluate*

Why a Surrogate

Since evaluating f is expensive, it should be done *infrequently*

The main trick to achieve this is using a *surrogate model*

- After each evaluation we train a *Machine Learning model*
- ...Then we perform optimization *on the ML model*
- ...Since it can be evaluate much more quickly

The process is usually start by sampling a few random points

This is where the name stems from


- Since we use the ML model instead of the function, we call it a *surrogate*

- Moreover, we optimize over *prior* information (i.e. the current model)
- ...And we refine the model based on the evaluation (*posterior*)
- Hence we call it *Bayesian Optimization*

A Running Example

Let's assume we want to minimize the following function over $[0, 1]$

```
In [2]: bbf = lambda x: (0.5 - x**2) * np.sin(2 * 4 * np.pi * x)
xrange = np.linspace(0, 1, 1000)
target = pd.Series(index=xrange, data=bbf(xrange))
util.plot_series(target, figsize=figsize)
```




- There multiple local minima, and the global minimum is at $\simeq 0.19$

A Running Example

Let's start by sampling a few points at random

```
In [3]: np.random.seed(42)
xtr = np.sort(np.random.random(4))
ytr = bbf(xtr)
util.plot_series(target, figsize=figsize)
plt.scatter(xtr, ytr, color='tab:orange');
```



- Using *only the orange points* we need to train a model
- ...That gives us a good idea for a new point to evaluate

Which properties should our surrogate model have?

Properties of a Good Surrogate

Our surrogate model should

Approximate very accurately all evaluated points

- Assuming the function is deterministic
- ...The available evaluations are exact values

Reflect our confidence level on unexplored regions

- If we have few samples in a certain region
- ...We might want to search there just to see what the function looks like

Can you think of a ML model with these properties?

Gaussian Process Surrogates

Gaussian Processes check all the boxes!

- They can interpolate very well known measurements
- They provide a confidence level that decays with distance from observations

Let's try to train a simple GP for our example

```
In [4]: kernel = RBF(0.01, (1e-3, 1e3)) + WhiteKernel(1e-3, (1e-6, 1e-2))
gpr = GaussianProcessRegressor(kernel=kernel, n_restarts_optimizer=9, normalizer=None)
gpr.fit(xtr.reshape(-1, 1), ytr);
gpr.kernel_
```

```
Out[4]: RBF(length_scale=0.0792) + WhiteKernel(noise_level=0.000126)
```


- We use an *RBF kernel* to capture the distance-based correlation
- We also use a *white noise kernel* to avoid numerical instability
- ...But we keep it at a low value since the target function is deterministic

Gaussian Process Surrogate

Let's inspect our Gaussian Process Surrogate

```
In [5]: pmean, pstd = gpr.predict(xrange.reshape(-1, 1), return_std=True)
util.plot_gp(target=target, figsize=figsize, pred=pd.Series(index=xrange, data=pmean, dtype=float))
```

```
std=pd.Series(index=xrange, data=pstd), samples=pd.Series(index
```




- All known points are interpolated (almost) exactly
- ...And the confidence intervals behave in an intuitive fashion

What to Optimize?

Now we need to search over the surrogate model

This is the same as choosing *which function to optimize*

```
In [6]: pmean, pstd = gpr.predict(xrange.reshape(-1, 1), return_std=True)
util.plot_gp(target=target, figsize=figsize, pred=pd.Series(index=xrange, da
std=pd.Series(index=xrange, data=pstd), samples=pd.Series(index
```



Which area does it make sense to explore, and why?

Acquisition Function

We need to account for both the *predictions* and their *confidence*

- Area with *low predictions* are promising
- ...But so are also areas with *high confidence*

This issue is solved in SBO by optimizig an *acquisition function*

...Which should balance *exploration* and *exploitation*.

- [Examples](#) include the Probability of improvement, the Expected Improvement
- ...And the Lower/Upper confidence bound

We will use the Lower Confidence Bound, which is given by:


$$LCB(x) = \mu(x) - Z_{\alpha}\sigma(x)$$

- Where $\mu(x)$ is the predicted mean, $\sigma(x)$ is the predicted standard deviation
- ...And Z_{α} is multiplier for a $\alpha\%$ Normal confidence interval

Lower Confidence Bound

Let's see an example in our case with $Z_{\alpha} = 2.5$

```
In [7]: lcb = pmean - 2.5 * pstd
util.plot_gp(target=target, figsize=figsize, pred=pd.Series(index=xrange, da
plt.plot(xrange, lcb, color='tab:red');
```




- We can then optimize via any method applicable to our surrogate
- E.g. [Nelder-Mead](#), Mathematical Programming, or even simple grid search

Lower Confidence Bound

Let's see which point we would choose in our case

```
In [8]: best_idx = np.argmin(lcb)
util.plot_gp(target=target, figsize=figsize, pred=pd.Series(index=xrange, da
plt.plot(xrange, lcb, color='tab:red');
plt.scatter(xrange[best_idx], lcb[best_idx], marker='x', color='tab:red');
```



- The x value with the best acquisition function is highlighted with a red "x"

Updating the Surrogate

Now we update our surrogate model

First, we evaluate f for the new point and grow our training set:

```
In [9]: xtr2 = np.hstack((xtr, [xrange[best_idx]]))
        ytr2 = np.hstack((ytr, [bbf(xrange[best_idx])]))
```

Then we can retrain our Gaussian Process:

```
In [10]: gpr2 = GaussianProcessRegressor(kernel=kernel, n_restarts_optimizer=9, normalize_y=True)
        gpr2.fit(xtr2.reshape(-1, 1), ytr2);
        gpr2.kernel_
```

```
Out[10]: RBF(length_scale=0.0999) + WhiteKernel(noise_level=2.46e-06)
```


- Then we should optimize the acquisition function again
- ...But we will limit ourselves to showing the updated predictions

Updating the Surrogate

Here are the estimates for the update surrogate

...Together with the acquisition function and the next iterate

```
In [11]: pmean2, pstd2 = gpr2.predict(xrange.reshape(-1, 1), return_std=True)
        lcb2 = pmean2 - 2.5 * pstd2
        best_idx2 = np.argmin(lcb2)
        util.plot_gp(target=target, figsize=figsize, pred=pd.Series(index=xrange, data=pmean2))
        plt.plot(xrange, lcb2, color='tab:red');
        plt.scatter(xrange[best_idx2], lcb2[best_idx2], marker='x', color='tab:red')
```



Surrogate-Based Bayesian Optimization

Let's review the general method

- Given a collection $\{x_i, y_i\}_i$ of evaluated points
- ...We train a surrogate-model $\hat{f}$ for f

Then we proceed as follows:

- We optimize an acquisition function $a_{\hat{f}}(x)$ to find a value x'
- We evaluate $y' = f(x')$
- If y' is better than the current optimum $f(x^*)$:
 - Then we replace x^* with x'
- We expand our collection of measurements to include (x', y')
- We retrain $\hat{f}$
- We repeat until a termination condition is reached

A Few Considerations

Different Bayesian optimization algorithms:

- Make use of different surrogate models
- Rely on different criteria for choosing x'
- Strike different trade-offs in terms of number of (expensive) evaluations of f
- ...And the quality of the obtained solutions

For more information, see (e.g.) [this tutorial](#)

In practice, you don't have to code from scratch

...Since multiple libraries are available, like:

- The `scikit-optimize` [package](#) (crude, reasonably fast, unfortunately unmaintained)
- The `bayesian-optimization` [python module](#) (more stable, but also slower)
- The [RBOpt solver](#) (stable, fast, more complex in terms of requirements)

SBO for Threshold Calibration

Back to Our Motivating Problem

We will use SBO to tackle our policy definition problem

$$\operatorname{argmin}_{\varepsilon} \sum_{k \in K} \operatorname{cost}(f(x_k, \theta^*), 1/2) \quad (1)$$

$$\text{s.t.: } \theta^* = \operatorname{argmin}_{\theta} L(f(x_k, \lambda), 1_{y_k \geq \varepsilon}) \quad (2)$$

Here's our plan:

- We need to optimize over ε
- Our goal is minimizing the cost
- Computing the cost requires to re-define the classes
- ...And therefore to repeat training

Our implementation will be based on [bayesian-optimization](#)

The Black Box Function

As a first step, we need to define our black box function

We will use a function class (in the `util` module) with this structure:

```
class ClassifierCost:
    def __init__(self, machines, X, y, cost_model, init_epochs=20,
inc_epochs=3):
        ...

    def __call__(self, params):
        ...
```

- In the constructor, we provide parameters that are fixed during optimization
- In the `__call__` method, we retrain the model and evaluate the cost
- The `__call__` method is executed when we try to invoke an object of this class
- ...Meaning that we can treat an object of this class as a normal function

The Black Box Function

It is worth having a deeper look at the `__call__` method

```
def __call__(self, params):
    theta = params[0] # There is only one parameter to optimize
```

```

        lbl = (self.y >= theta) # Redefine classes
        # Determine the number of epochs and retrain
        epochs = self.init_epochs if not self.is_init else
self.inc_epochs
        self.is_init = True
        train_nn_model(self.nn, self.X, lbl,
loss='binary_crossentropy', epochs=epochs,
                        verbose=0, patience=10, batch_size=32,
validation_split=0.2)
        ...

```

- At each execution we redefine the classes
- We use warm starting to make the process faster
- Each training attempt after the first uses only a few epochs

The Black Box Function

It is worth having a deeper look at the `__call__` method

```

def __call__(self, params):
    ...
    self.stored_weights[theta] = self.nn.get_weights() # Store weights
    # Evaluate cost
    pred = np.round(self.nn.predict(self.X, verbose=0).ravel())
    cost, fails, slack = self.cost_model.cost(self.machines, pred,
0.5, return_margin=True)
    return cost

```

- We store the weights in a dictionary for later retrieval
- We need this to rebuild the optimal network once optimization is over
- Finally, we evaluate the cost
- The actual code in `util` also prints some information

The Black Box Function

We can build an object in the usual way

```
In [12]: ccf = util.ClassifierCost(machines=tr['machine'], X=tr_s[dt_in], y=tr['rul'])
```

...But since it is a function, we can *invoke* it:

```
In [13]: ccf(20, verbose=True)
```

```
eps: 20.00, avg. cost: -86.67, avg. fails: 0.00, avg. slack: 29.82
```

```
Out[13]: np.int64(16120)
```

- The function returns the negated cost
- ...Since `bayesian-optimization` is designed for maximization

Running the Solver

Now we can define our box constraints and run the optimization process

```
In [14]: pbounds = {'eps': (1, 20)} # Box constraints
optimizer = BayesianOptimization(f=ccf, pbounds=pbounds, random_state=42)
optimizer.maximize(init_points=3, n_iter=10)
```

iter	target	eps
1	18606.0	8.1162622
2	15848.0	19.063571
3	16821.0	14.907884
4	-54286.0	1.1520215
5	17379.0	11.328003
6	16962.0	17.058181
7	18255.0	9.5110781
8	17478.0	13.096088
9	16158.0	20.0
10	18819.0	8.6924914
11	18288.0	8.4601542
12	17478.0	5.6899704
13	18069.0	6.7183624

- The implementation is very close to what we have showed
- The results will be a bit noisy, since training is stochastic

Retrieve the Results

We can access the best ϵ value from a result data structure

```
In [15]: print(optimizer.max)
best_eps = optimizer.max['params']['eps']

{'target': np.float64(18819.0), 'params': {'eps': np.float64(8.692491437271586)}}
```

We will use it to retrieve the weights of the best network:

```
In [16]: nn = keras.models.clone_model(ccf.nn)
nn.set_weights(ccf.stored_weights[best_eps])
```

- This is possible since we stored all the tested weights in our class

Evaluate the Classifier

Finally, we can evaluate our classifier

```
In [17]: tr_pred = np.round(nn.predict(tr_s[dt_in], verbose=0).ravel())
         ts_pred = np.round(nn.predict(ts_s[dt_in], verbose=0).ravel())

         tr_c, tr_f, tr_sl = cmodel.cost(tr['machine'].values, tr_pred, 0.5, return_n
         ts_c, ts_f, ts_sl = cmodel.cost(ts['machine'].values, ts_pred, 0.5, return_n
         print(f'Avg. cost: {tr_c/len(tr_mcn)} (training), {ts_c//len(ts_mcn)} (test)
         print(f'Avg. fails: {tr_f/len(tr_mcn)} (training), {ts_f/len(ts_mcn)} (test)
         print(f'Avg. slack: {tr_sl/len(tr_mcn):.2f} (training), {ts_sl/len(ts_mcn):.
```

```
Avg. cost: -101.1774193548387 (training), -112 (test)
Avg. fails: 0.0 (training), 0.0 (test)
Avg. slack: 15.03 (training), 11.98 (test)
```

- The performance is now better than the regression approach!

AutoML

Many ML models have hyper parameters!

...And tuning them may sometimes improve the performance

- The problem is that tuning multiple parameters may be complicated
- ...And every training attempts is expensive

This makes hyper-parameter tuning a perfect application for SBO

...And other similar approaches. A few libraries you might have heard of:

- [Hyperopt](#)
- [Optuna](#)

In recent years the concept has been generalized to AutoML

...Where we can start chanking the architecture and model type, too!

- It's a big topic (and big techs have some available SW solutions)
- A good starting reference is [this web site](#)